

# King County House Sales Analysis

---

Author: [Femi Kamau \(https://www.github.com/ctrl-Karugu\)](https://www.github.com/ctrl-Karugu)

## Overview

In our analysis, we explored the data provided by the stakeholder and build a multiple linear features stipulated in the dataset. From there, we analysed the results and came to a conclusion that have a significant impact on the price of a house in King County:

- Have a house by the water
- Increase the number of bedrooms
- Improve the overall grade of the home
- Increase the number of floors
- Increase the size of the basement
- Strive to maintain the house to ensure that it is in good condition

## 1. Business Understanding

A real estate agency located in King County is looking to advice home owners about how to increase the value of their homes. The agency is looking to use the King County house data provided to make renovations to make to increase the value of a home.

## 2. Data Understanding

This phase is broken down into four tasks together with its projected outcome or output:

- Collect Initial Data
- Describe Data
- Explore Data
- Verify Data Quality

There was no need to collect any data for this project as it was already provided by the stakeholder as house data from King County and is in .csv format.

Load Libraries

In [274]:

```
# data
import numpy as np
import pandas as pd

# visualization
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno
import folium
import warnings

# modeling
import statsmodels.api as sm
from sklearn.metrics import mean_absolute_error
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.tools.tools import add_constant

# statistics
import scipy.stats as stats

# styling
plt.style.use('seaborn')
sns.set_style('whitegrid')

warnings.filterwarnings('ignore')
```

Import Data

```
In [275]:
```

```
# King County House Sales dataset is imported and assigned to the variable 'data'
data = pd.read_csv('../data/raw/kc_house_data.csv')

# The shape of the dataframe and the last 5 rows are outputted
print(data.shape)
data.tail()
```

```
(21597, 21)
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
<b>21592</b>	263000018	5/21/2014	360000.0	3	2.50	1530	1131	3.0	NO
<b>21593</b>	6600060120	2/23/2015	400000.0	4	2.50	2310	5813	2.0	NO
<b>21594</b>	1523300141	6/23/2014	402101.0	2	0.75	1020	1350	2.0	NO
<b>21595</b>	291310100	1/16/2015	400000.0	3	2.50	1600	2388	2.0	NaN
<b>21596</b>	1523300157	10/15/2014	325000.0	2	0.75	1020	1076	2.0	NO

```
5 rows × 21 columns
```

There are 21 columns and 21597 rows in the dataset:

- **Numerical Columns (15)**

date - Date house was sold

price - Sale price (prediction target)

bedrooms - Number of bedrooms

bathrooms - Number of bathrooms

sqft\_living - Square footage of living space in the home

sqft\_lot - Square footage of the lot

floors - Number of floors (levels) in house

sqft\_above - Square footage of house apart from basement

sqft\_basement - Square footage of the basement

yr\_built - Year when house was built

yr\_renovated - Year when house was renovated

lat - Latitude coordinate

long - Longitude coordinate

sqft\_living15 - The square footage of interior housing living space for the nearest 15 n

sqft\_lot15 - The square footage of the land lots of the nearest 15 neighbors

- **Categorical Columns (6)**

id - Unique ID for each home sold

waterfront - Whether the house has a view to a waterfront

view - An index from 0 to 4 of how good the view of the property was

condition - An index from 1 to 5 on the condition of the house

grade - An index from 1 to 13, where 1-3 falls short of building construction and design, construction and design, and 11-13 have a high quality level of construction and design

zipcode - What zipcode area the house is in

```
In [276]:
```


```
# Describe the data
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               21597 non-null  int64
1   date            21597 non-null  object
2   price           21597 non-null  float64
3   bedrooms        21597 non-null  int64
4   bathrooms       21597 non-null  float64
5   sqft_living     21597 non-null  int64
6   sqft_lot        21597 non-null  int64
7   floors          21597 non-null  float64
8   waterfront      19221 non-null  object
9   view            21534 non-null  object
10  condition       21597 non-null  object
11  grade           21597 non-null  object
12  sqft_above      21597 non-null  int64
13  sqft_basement   21597 non-null  object
14  yr_built        21597 non-null  int64
15  yr_renovated    17755 non-null  float64
16  zipcode         21597 non-null  int64
17  lat             21597 non-null  float64
18  long            21597 non-null  float64
19  sqft_living15   21597 non-null  int64
20  sqft_lot15      21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB
```

```
In [277]:
```

```
# Visualise the missing values in the dataset
msno.bar(data, color='lightgreen', figsize=(10, 5), fontsize=8)
plt.title('Missing Values Within Dataset');
```



From the barplot above, we can see that the columns with missing data are `waterfront`, `view`, and `yr_renovated`. The `waterfront` column has 2376 missing values, the `view` column has 63 missing values, and the `yr_renovated` column has 3842 missing values. This accounts for 11%, 0.3%, and 18% of the total number of rows in the dataset.

## 2.1 Univariate Analysis

In this section, we'll explore each column in the dataset to see the distributions of features and gain some useful insights. The main two parts in this section are:

- Categorical Columns
- Numerical Columns

### 2.1.1 Categorical Columns

There are 5 Categorical Columns in the dataset that we shall be analysing:

- id
- waterfront
- view
- condition
- grade
- zipcode

Functions to visualise the data in the categorical columns

In [278]:

```
# Fuction to get the value counts of the data in the columns
def get_value_counts(df, col):
    ''' Returns the value counts of a column in a dataframe '''
    counts = df[col].value_counts(dropna=False)
    return counts

# Function to visualise the the data in the columns
def plot_data(df, col, title):
    ''' Plots the value counts of a column in a dataframe as a bar chart '''
    get_value_counts(df, col).plot(kind='bar', figsize=(10, 5), color='lightgreen', edgecolor='black')
    plt.title(title)
    plt.xticks(rotation=0);
```

#### 2.1.1.1 ID

---

The `id` column is a unique identifier for each house sold.

---

The univariate analysis of the `id` column will be less about identifying the data distribution, but more about the number of unique values in the column. From the count of the unique values we will be able to identify any duplicates.

In [279]:

```
# Check for duplicates in the 'id' column
data.id.duplicated().sum()
```

We see that there are 177 duplicated ids in the dataset. This could mean that there are some more than once, or it could also mean that there are some records that have been imputed. We will have to investigate this further in the data preparation phase.

### 2.1.1.2 Waterfront


The `waterfront` column identifies whether the house is on a waterfront or not.

In [280]:

```
# Identify the unique values (and counts) in the 'waterfront' column
print(get_value_counts(data, 'waterfront'))

# Visualise the data distribution
plot_data(data, 'waterfront', 'Waterfront Column Data Distribution')
```

```
NO      19075
NaN      2376
YES       146
Name: waterfront, dtype: int64
```



The distribution above shows that most of the houses in the dataset are not on a waterfront. The count for 'waterfront' is 146, which is 0.7% of the total number of houses in the dataset. The missing values (nan) represent 2,376 houses, which is approximately 12.2% of the total.

which is 11% of the total number of rows in the dataset. As this is a categorical column, we will use the mode of the column.

### 2.1.1.3 View


The view column identifies the quality of view from the house.

In [281]:

```
# Identify the unique values (and counts) in the 'view' column
print(get_value_counts(data, 'view'))

# Visualise the data distribution
plot_data(data, 'view', 'View Column Data Distribution')
```

```
NONE          19422
AVERAGE       957
GOOD           508
FAIR           330
EXCELLENT      317
NaN            63
Name: view, dtype: int64
```



In the distribution above, we see that majority of the houses in the dataset have a no view. F



in this columns are 63, which is 0.29% of the total number of rows in the dataset. As this is a number of rows in the dataset, we can drop the rows with missing values in this column.

### 2.1.1.4 Condition


The condition column identifies the condition of the house.

In [282]:

```
# Identify the unique values (and counts) in the 'condition' column
print(get_value_counts(data, 'condition'))

# Visualise the data distribution
plot_data(data, 'condition', 'Condition Column Data Distribution')
```

```
Average      14020
Good          5677
Very Good     1701
Fair           170
Poor           29
Name: condition, dtype: int64
```



From the distribution above, we can see that most of the houses in the dataset are in average houses in average condition is 12437, this accounts for 57.6% of the total number of houses

houses in good condition are 5041, this accounts for 23.3% of the total number of houses in the dataset. Houses in very good condition are 1509, this accounts for 7% of the total number of houses in the dataset. Houses in fair condition are 152, this accounts for 0.7% of the total number of houses in the dataset. Houses in poor condition are 25, this accounts for 0.1% of the total number of houses in the dataset. Furthermore, there is no missing data within this column.

#### 2.1.1.5 Grade


The grade column identifies the quality of construction and design of the house. The grade column also identifies the construction quality of improvements. Grades run from grade 1 to 13.

```
In [283]:
```

```
# Identify the unique values (and counts) in the 'grade' column
print(get_value_counts(data, 'grade'))

# Visualise the data distribution
plot_data(data, 'grade', 'Grade Column Data Distribution')
plt.xticks(rotation=45);
```

```
7 Average      8974
8 Good         6065
9 Better       2615
6 Low Average  2038
10 Very Good   1134
11 Excellent    399
5 Fair         242
12 Luxury       89
4 Low          27
13 Mansion     13
3 Poor          1
Name: grade, dtype: int64
```



From the distribution above, we see that the houses in this dataset range from grades 3-13.

evenly distributed as we can see majority of the houses with a grade of 7 (representing Average Good). Lastly, there are no missing values within this column.

#### 2.1.1.5 Zipcode


The zipcode column identifies the zipcode area the house is in.

```
In [284]:
```

```
# Identify the unique values (and counts) in the 'zipcode' column
print(get_value_counts(data, 'zipcode'))

# Visualise the data distribution
plot_data(data, 'zipcode', 'Zipcode Column Data Distribution')
plt.xticks(rotation=90);
```

```
98103    602
98038    589
98115    583
98052    574
98117    553
...
98102    104
98010    100
98024     80
98148     57
98039     50
Name: zipcode, Length: 70, dtype: int64
```



From the distribution above, we see that the zipcode with the most houses is 98103. The zip 98039. Unlike the other categorical columns, we see more evenly distributed data in this col

this column

## Summary Of The Categorical Columns

- The quality of the data in the categorical columns is fairly good. Other than a few missing view columns, and duplicated values in the id column, the data is good to work with.

### 2.1.2 Numerical Columns

There are 15 Numerical Columns in the dataset that we shall be analysing:

- date
- price
- bedrooms
- bathrooms
- sqft\_living
- sqft\_lot
- floors
- sqft\_above
- sqft\_basement
- yr\_built
- yr\_renovated
- lat
- long
- sqft\_living15
- sqft\_lot15

Functions to visualise the data in the numerical columns

In [285]:

```
# Function that describes the statistics of the data
def describe_data(df, col):
    ''' Returns the statistics of a column in a dataframe '''
    print(df[col].describe())

# Function to plot the histogram, kde and boxplot of the data
def plot_distribution(df, col, title, bins_=10):
    ''' Plots the distribution of a column in a dataframe as a histogram, kde and boxplot.
    # creating a figure composed of two matplotlib.Axes objects (ax_box and ax_hist)
    f, (ax_box, ax_hist) = plt.subplots(2, sharex=True, gridspec_kw={"height_ratios": (1, 1)})

    # assign a graph to each ax
    sns.boxplot(df[col], ax=ax_box, color='lightgreen')
    sns.histplot(data=df, x=col, ax=ax_hist, kde=True, color='lightgreen', bins=bins_,
    plt.suptitle(title)
    plt.tight_layout();
```

### 2.1.2.1 Date

The `date` column identifies the date the house was sold.

Date can either be a categorical or numerical column. In this case, we will treat it as a numerical column.

```
In [286]:
```

```
# Print all the unique values in the 'date' column
```

```
print(get_value_counts(data, 'date').index.tolist())
```

```
# Get the unique values (and counts) in the 'date' column
```

```
print(get_value_counts(data, 'date'))
```

```
['6/23/2014', '6/25/2014', '6/26/2014', '7/8/2014', '4/27/2015', '3/25/2015', '7/9/2014', '4/14/2015', '4/28/2014', '4/21/2015', '8/26/2014', '10/28/2014', '7/14/2014', '5/20/2014', '7/1/2014', '8/20/2014', '6/17/2014', '4/8/2015', '7/16/2014', '5/28/2014', '4/23/2015', '8/5/2014', '3/27/2015', '8/25/2014', '6/20/2014', '7/23/2014', '3/26/2015', '6/3/2014', '5/27/2014', '8/22/2014', '9/23/2014', '4/2/2015', '4/24/2015', '7/25/2014', '4/7/2015', '8/27/2014', '6/19/2014', '3/24/2015', '6/4/2014', '11/13/2014', '8/12/2014', '3/4/2015', '7/18/2014', '9/24/2015', '6/10/2014', '7/21/2014', '6/16/2014', '12/2/2014', '8/13/2014', '10/27/2014', '12/1/2014', '2/25/2015', '9/9/2014', '7/28/2014', '5/21/2014', '5/5/2015', '7/24/2014', '9/9/2014', '5/7/2014', '8/19/2014', '9/16/2014', '3/23/2015', '8/14/2014', '10/29/2014', '4/9/2015', '5/22/2014', '10/7/2014', '9/22/2014', '3/30/2015', '8/4/2014', '7/2/2014', '7/10/2014', '9/26/2014', '10/21/2014', '6/12/2014', '10/14/2014', '4/13/2015', '5/6/2015', '9/1/2015', '11/18/2014', '8/21/2014', '11/20/2014', '5/13/2014', '9/10/2014', '9/5/2014', '7/17/2014', '10/1/2014', '10/15/2014', '8/6/2014', '9/3/2014', '10/30/2014', '8/18/2014', '5/5/2014', '4/6/2015', '10/20/2014', '5/2/2015', '9/29/2014', '5/6/2014', '5/19/2014', '4/17/2015', '4/30/2015', '7/31/2014', '11/17/2014', '5/15/2014', '11/19/2014', '11/21/2014', '5/9/2014', '8/1/2014', '11/10/2014', '5/14/2014', '5/8/2014', '9/2/2014', '5/12/2014', '12/15/2014', '10/9/2014', '9/4/2014', '11/24/2014', '9/18/2014', '9/25/2014', '6/13/2014', '3/5/2015', '12/28/2014', '6/2/2014', '8/8/2014', '4/16/2015', '9/15/2014', '4/10/2015', '6/30/2014', '5/1/2015', '12/11/2014', '11/5/2014', '5/7/2015', '9/19/2014', '3/12/2015', '12/10/2014', '7/30/2014', '4/3/2015', '2/19/2015', '10/8/2014', '10/17/2014', '12/8/2014', '9/8/2014', '9/17/2014', '3/10/2015', '7/3/2014', '11/7/2014', '10/2/2014', '3/3/2014', '2/20/2015', '11/4/2014', '2/23/2015', '11/25/2014', '3/13/2015', '3/19/2015', '10/6/2014', '2/13/2014', '12/12/2014', '3/9/2015', '2/17/2015', '5/2/2014', '12/18/2014', '10/3/2014', '12/22/2014', '6/6/2014', '1/29/2014', '12/17/2014', '3/20/2015', '10/13/2014', '3/3/2015', '8/7/2014', '12/23/2014', '2/11/2015', '8/15/2014', '1/28/2015', '2/26/2015', '1/5/2015', '10/24/2014', '3/6/2015', '8/29/2014', '1/27/2015', '2/9/2015', '1/2/2015', '2/10/2015', '1/7/2015', '1/16/2015', '2/4/2015', '1/14/2015', '2/6/2015', '2/5/2015', '2/27/2015', '11/2015', '1/22/2015', '2/12/2015', '1/26/2015', '1/23/2015', '2/2/2015', '11/26/2014', '2/3/2015', '1/20/2015', '1/15/2015', '1/29/2015', '1/13/2015', '3/2/2015', '12/30/2014', '1/12/2015', '1/6/2015', '12/19/2014', '5/11/2014', '1/9/2015', '1/30/2015', '12/24/2014', '5/13/2015', '4/25/2015', '3/21/2015', '4/26/2015', '4/12/2015', '2/22/2015', '6/22/2014', '5/14/2015', '5/24/2014', '6/8/2014', '5/3/2015', '7/12/2014', '1/19/2015', '3/29/2014', '6/21/2014', '6/14/2014', '3/28/2015', '7/20/2014', '7/26/2014', '8/23/2014', '6/29/2014', '6/15/2014', '5/1/2014', '4/5/2015', '7/5/2014', '2/16/2015', '3/1/2015', '5/31/2014', '5/2/2015', '4/19/2015', '10/18/2014', '14/2015', '11/1/2014', '9/13/2014', '5/4/2014', '10/25/2014', '5/25/2014', '9/21/2014', '5/10/2014', '9/6/2014', '9/27/2014', '4/18/2015', '2/28/2015', '7/19/2014', '11/8/2014', '3/22/2015', '5/3/2014', '8/17/2014', '12/13/2014', '4/4/2015', '10/19/2014', '6/7/2014', '12/14/2014', '8/16/2014', '9/14/2014', '11/23/2014', '10/26/2014', '1/25/2015', '9/28/2014', '5/9/2015', '8/10/2014', '11/29/2014', '11/16/2014', '2/14/2015', '10/12/2014', '2/7/2014', '8/31/2014', '10/5/2014', '7/6/2014', '2/21/2015', '7/4/2014', '8/24/2014', '2/1/2015', '10/11/2014', '1/27/2014', '5/11/2014', '12/21/2014', '8/9/2014', '9/7/2014', '11/15/2014', '11/28/2014', '1/10/2015', '5/27/2014', '2/15/2015', '3/8/2015', '8/30/2014', '5/15/2015', '1/17/2015', '11/2/2014', '1/31/2015', '5/24/2015', '5/6/23/2014', 142
6/25/2014    131
6/26/2014    131
7/8/2014     127
4/27/2015    126
...
11/2/2014     1
1/31/2015     1
5/24/2015     1
5/17/2014     1
7/27/2014     1
Name: date, Length: 372, dtype: int64
```



From the output above, we can see that the data has been stored in string format. We will have the data in datetime format in the data preparation phase. Furthermore, it seems that most of the houses were built in 2015.

#### 2.1.2.2 Price

---

The price column identifies the price of the house.

---

```
In [287]:
```

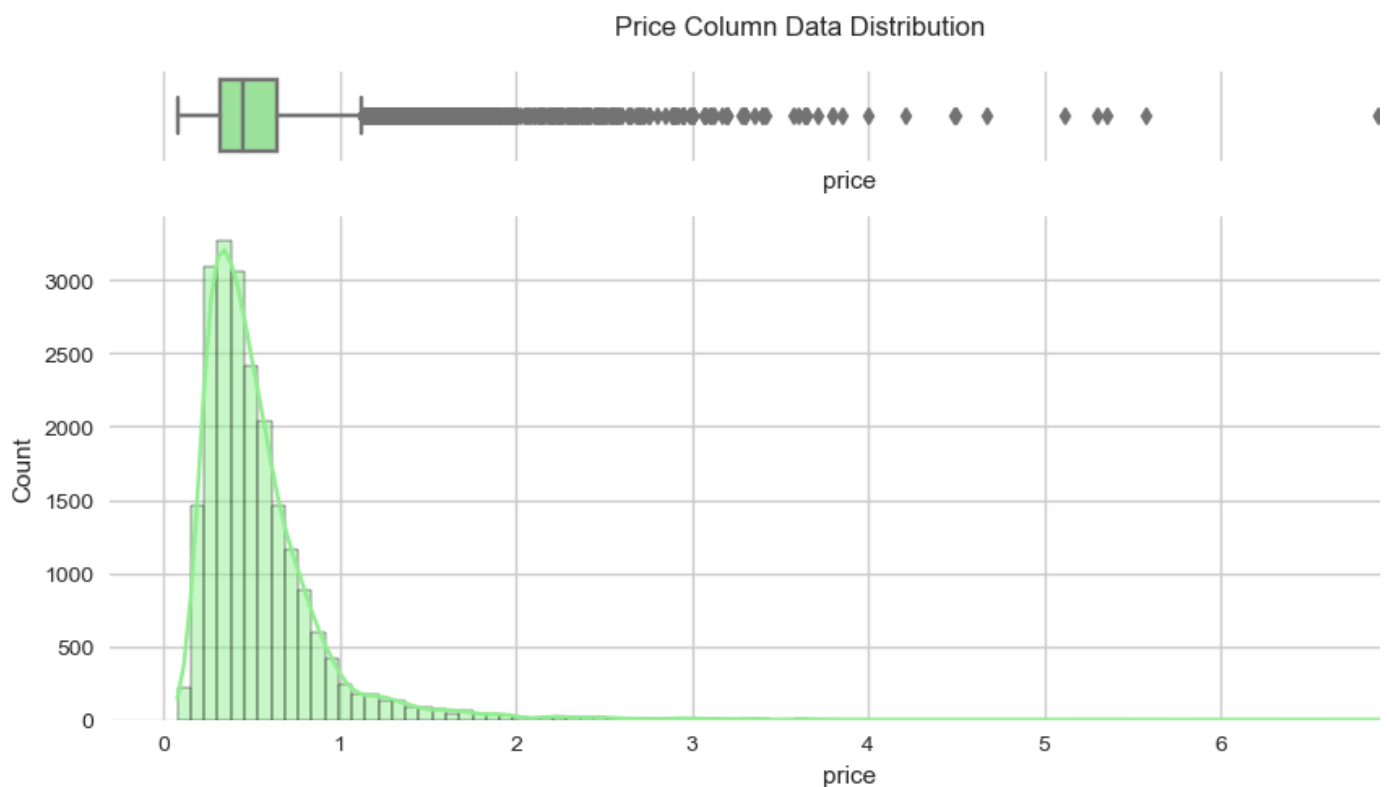
```
# Describe the 'price' column
```

```
describe_data(data, 'price')
```

```
# Visualise the data distribution
```

```
plot_distribution(data, 'price', 'Price Column Data Distribution', 100)
```

```
count    2.159700e+04
mean      5.402966e+05
std       3.673681e+05
min       7.800000e+04
25%       3.220000e+05
50%       4.500000e+05
75%       6.450000e+05
max       7.700000e+06
Name: price, dtype: float64
```



From the distribution above, we see that the price column is skewed to the right. This means homes in the dataset are . The minimum price of a house in the dataset is 78,000, and the maximum price in the dataset is 7,700,000. The mean price of a house in the dataset is 540,297, and the median price is 450,000. The standard deviation of the price column is 367,368.

Looking at the kurtosis of the distribution shows that

### 2.1.2.3 Bedrooms

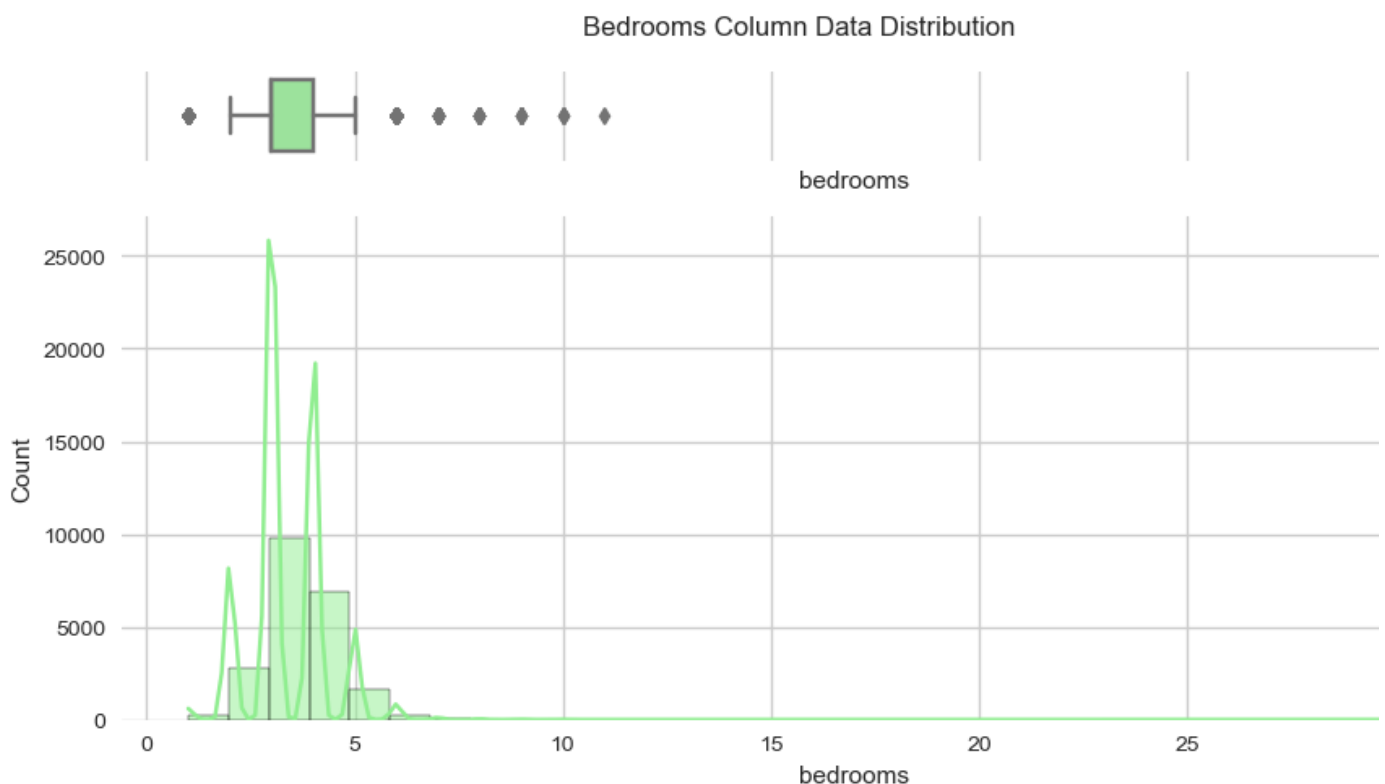
The bedrooms column identifies the number of bedrooms in the house.

In [288]:

```
# Describe the 'bedroom' column
describe_data(data, 'bedrooms')

# Visualise the data distribution
plot_distribution(data, 'bedrooms', 'Bedrooms Column Data Distribution', 33)
```

```
count    21597.000000
mean       3.373200
std        0.926299
min        1.000000
25%        3.000000
50%        3.000000
75%        4.000000
max       33.000000
Name: bedrooms, dtype: float64
```



The bedroom column distribution is not skewed as the and is normally distributed. The minimum house in the dataset is 1, and the maximum number of bedrooms in a house in the dataset is 3.37, and the median number of bedrooms in a house standard deviation of the bedrooms column is 0.93.

### 2.1.2.4 Bathrooms

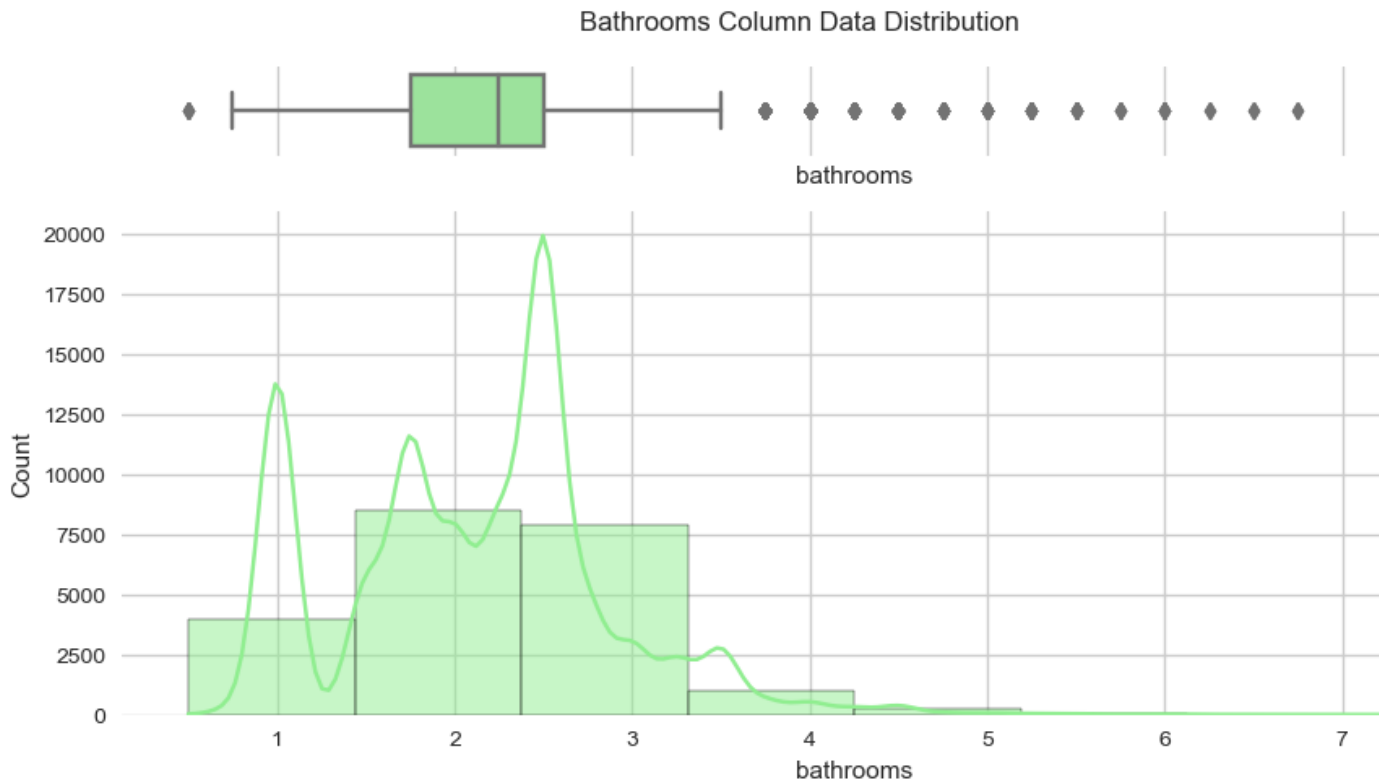
The bathrooms column identifies the number of bathrooms in the house.

In [289]:

```
# Describe the 'bathrooms' column
describe_data(data, 'bathrooms')

# Visualise the data distribution
plot_distribution(data, 'bathrooms', 'Bathrooms Column Data Distribution', 8)
```

```
count    21597.000000
mean       2.115826
std        0.768984
min        0.500000
25%        1.750000
50%        2.250000
75%        2.500000
max        8.000000
Name: bathrooms, dtype: float64
```



From the distribution above we can see that the bathroom column is not skewed. This is because the distribution is almost symmetric. The minimum number of bathrooms in a house in the dataset is 0.5, and the maximum number of bathrooms in a house in the dataset is 8. The mean number of bathrooms in a house in the dataset is 2.115826. The standard deviation of the bathrooms column is 0.768984. The number of bathrooms in a house in the dataset is 2.25. The standard deviation of the bathrooms column is 0.768984.

### 2.1.2.5 Sqft Living

The sqft living column identifies the square footage of the house.

In [290]:

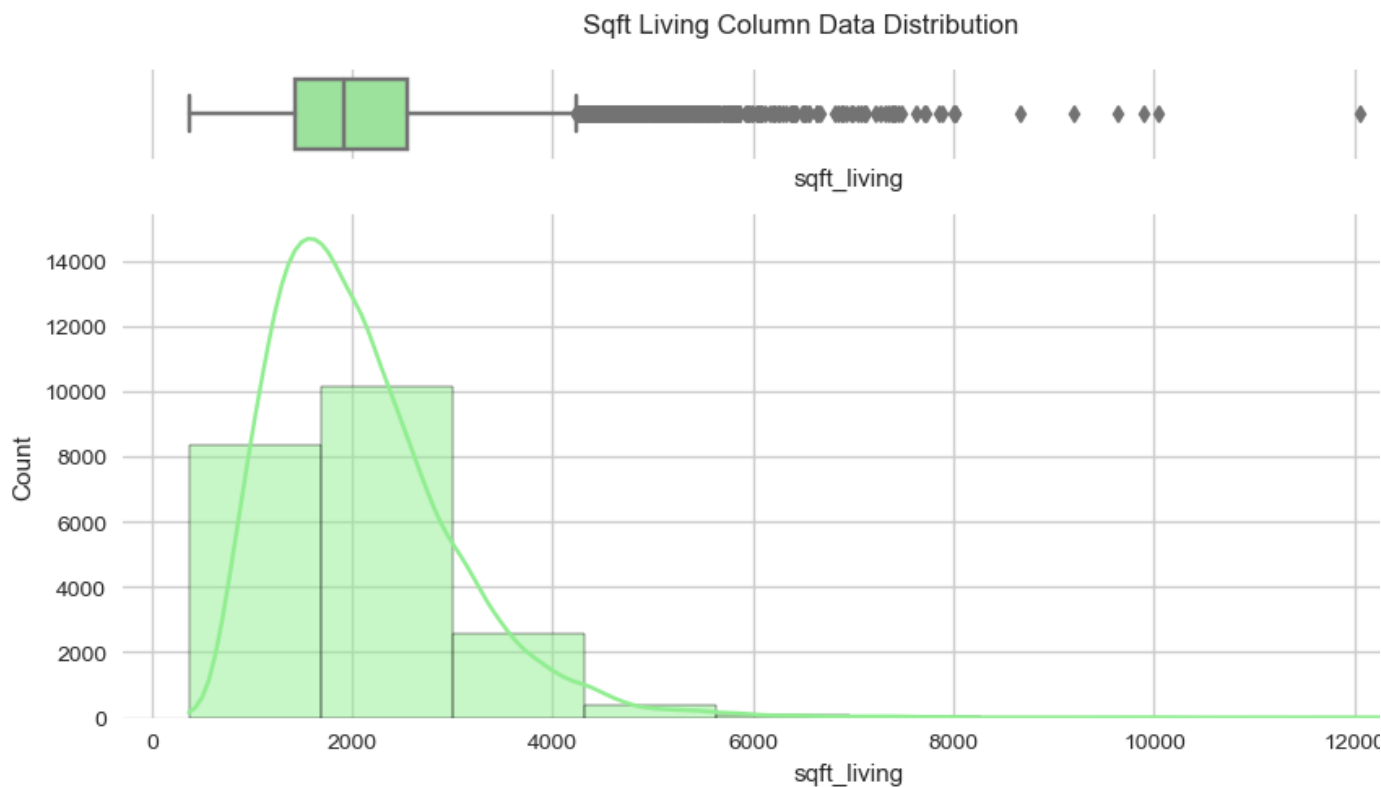
```
# Describe the 'sqft_living' column
```

```
describe_data(data, 'sqft_living')
```

```
# Visualise the data distribution
```

```
plot_distribution(data, 'sqft_living', 'Sqft Living Column Data Distribution')
```

```
count    21597.000000
mean      2080.321850
std       918.106125
min       370.000000
25%      1430.000000
50%      1910.000000
75%      2550.000000
max      13540.000000
Name: sqft_living, dtype: float64
```



From the distribution above, we can see that the sqft living column is skewed to the right. The footprint of the homes is greater than the median. The minimum square footage of a house in the dataset is 370. The maximum square footage of a house in the dataset is 13,540. The mean square footage of a

and the median square footage of a house in the dataset is 1910. The standard deviation of

### 2.1.2.6 Sqft Lot


The `sqft_lot` column identifies the square footage of the lot.

In [291]:

```
# Describe the 'sqft_lot' column
describe_data(data, 'sqft_lot')

# Visualise the data distribution
plot_distribution(data, 'sqft_lot', 'Sqft Lot Column Data Distribution', 100)
```

```
count    2.159700e+04
mean      1.509941e+04
std       4.141264e+04
min       5.200000e+02
25%       5.040000e+03
50%       7.618000e+03
75%      1.068500e+04
max      1.651359e+06
Name: sqft_lot, dtype: float64
```



From the distribution above, we can see that the data is skewed to the right. This is because

median. The minimum lot square footage is 520, the maximum lot square footage is 1,651,3  
feetage is 15,000, and the median lot square footage is 7618. The standard deviation of the

### 2.1.2.7 Floors

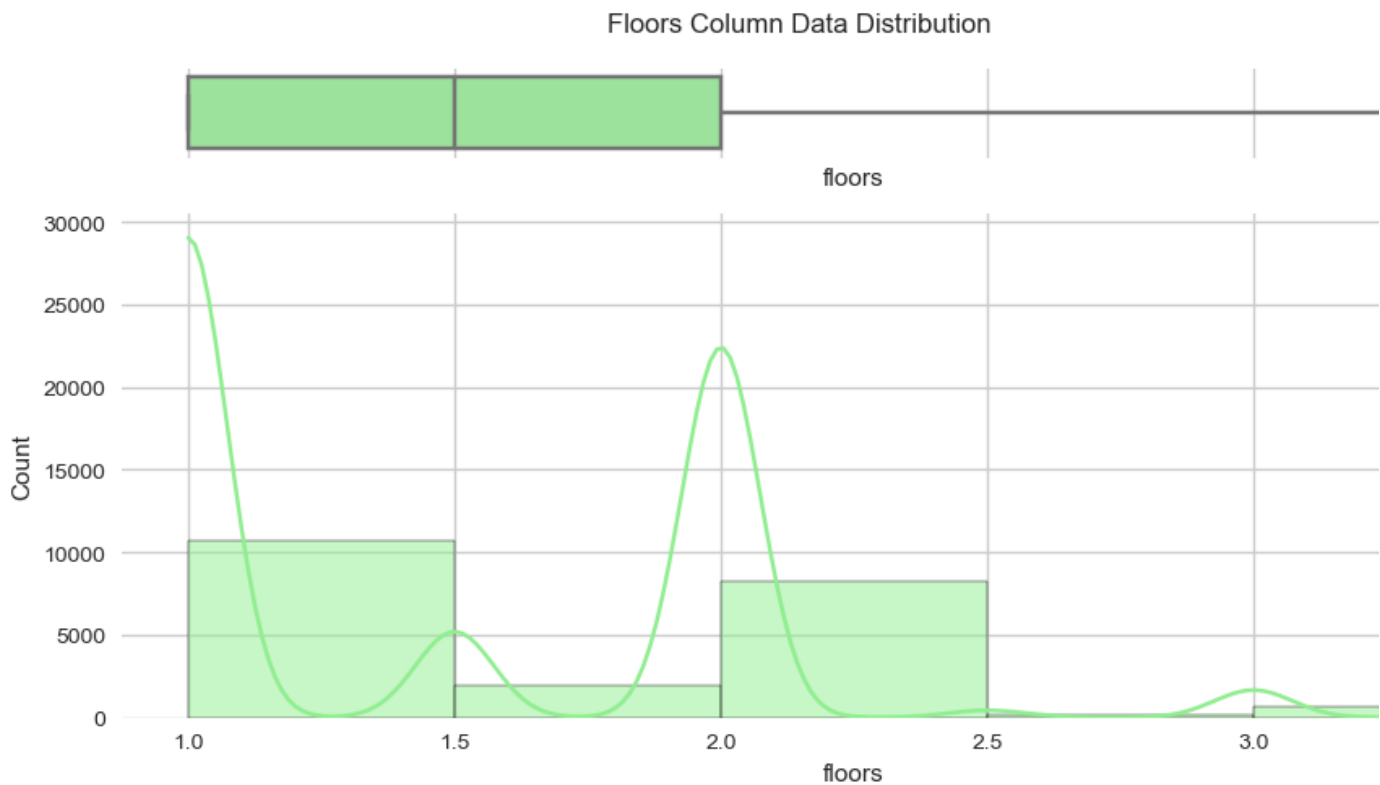
floors column identifies the number of floors in the house.

In [292]:

```
# Describe the 'floors' column
describe_data(data, 'floors')

# Visualise the data distribution
plot_distribution(data, 'floors', 'Floors Column Data Distribution', 5)
```

count	21597.000000
mean	1.494096
std	0.539683
min	1.000000
25%	1.000000
50%	1.500000
75%	2.000000
max	3.500000
Name: floors, dtype: float64	



From the distributions above, there is no particular trend in the floors column data. Majority c

1 floors. The minimum number of floors in a house is 1, and the maximum number of floors in this dataset is 1.5, and the mean number of floors in this dataset is approximately 1.25, and the standard deviation of the floors column is 0.54.

### 2.1.2.8 Sqft Above

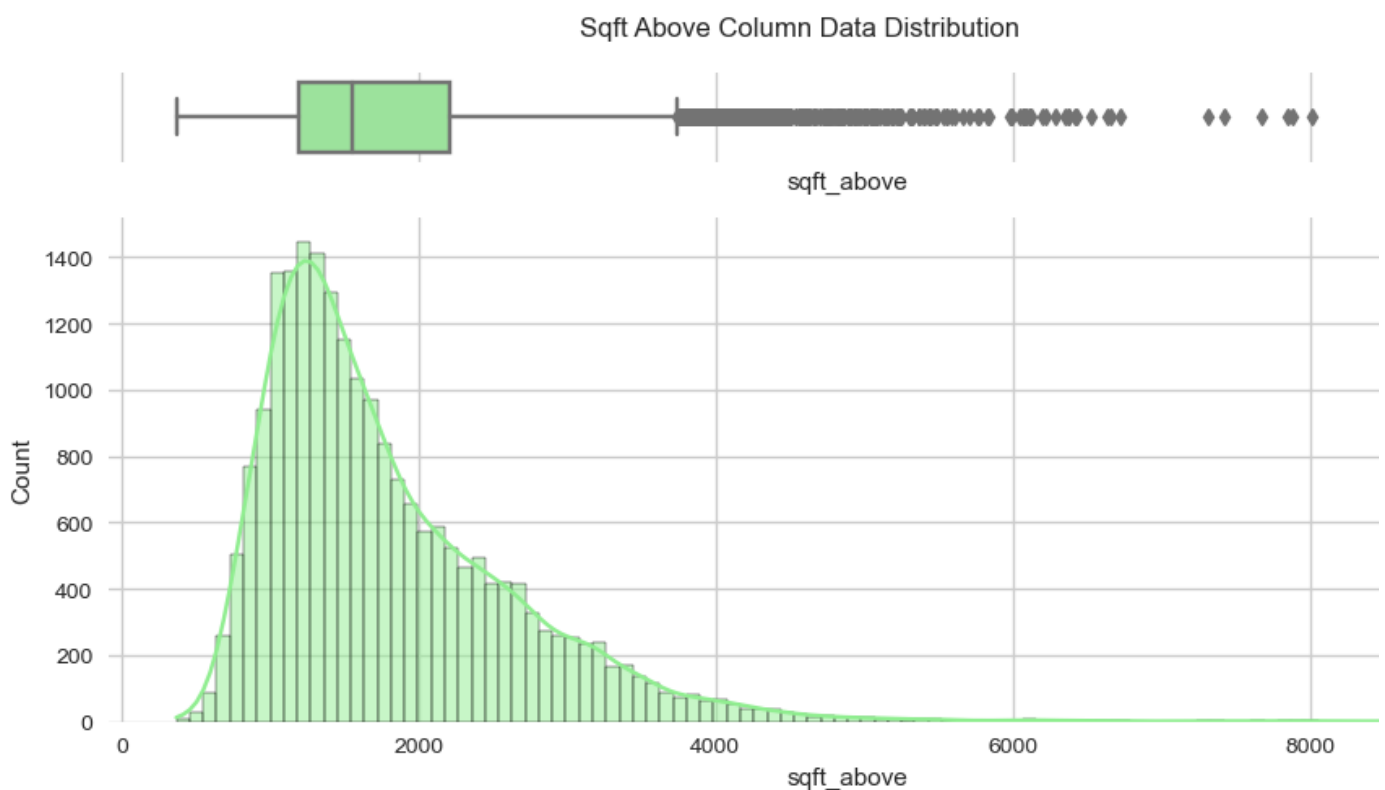
The `sqft_above` column identifies the square footage of the house above the ground.

In [293]:

```
# Describe the 'sqft_above' column
describe_data(data, 'sqft_above')

# Visualise the data distribution
plot_distribution(data, 'sqft_above', 'Sqft Above Column Data Distribution', 100)
```

```
count    21597.000000
mean      1788.596842
std       827.759761
min       370.000000
25%      1190.000000
50%      1560.000000
75%      2210.000000
max       9410.000000
Name: sqft_above, dtype: float64
```





From the distributions above, we see that the square footage above ground of the houses is right-skewed. This is because the mean is greater than the median. The minimum square footage above ground is 9,410. The maximum square footage of a house above ground is 9,410. The mean square footage above ground is 1,560. The median square footage above ground is 1,560. The standard deviation of the sqft above column is 1,560.

#### 2.1.2.9 Sqft Basement

The `sqft_basement` column identifies the square footage of the basement of the houses.

As this column is of the type `object`, we cannot do a distribution like the other numerical columns. We can view the contents of the column using the same technique as the categorical columns.

In [294]:

```
# Print all the unique values in the 'sqft_basement' column
print(get_value_counts(data, 'sqft_basement').index.tolist())

# Get the unique values (and counts) in the 'sqft_basement' column
print(get_value_counts(data, 'sqft_basement'))
```

```
['0.0', '?', '600.0', '500.0', '700.0', '800.0', '400.0', '1000.0', '900.0', '300.0', '200.0', '750.0', '450.0', '620.0', '580.0', '840.0', '420.0', '860.0', '1100.0', '670.0', '780.0', '550.0', '650.0', '240.0', '380.0', '770.0', '940.0', '910.0', '440.0', '880.0', '290.0', '1200.0', '350.0', '520.0', '920.0', '630.0', '730.0', '1010.0', '760.0', '640.0', '280.0', '340.0', '950.0', '820.0', '570.0', '560.0', '460.0', '790.0', '1060.0', '540.0', '810.0', '1040.0', '250.0', '140.0', '120.0', '890.0', '990.0', '1020.0', '470.0', '1070.0', '1250.0', '330.0', '390.0', '690.0', '610.0', '1030.0', '270.0', '150.0', '970.0', '1120.0', '220.0', '100.0', '260.0', '1050.0', '1300.0', '320.0', '710.0', '1400.0', '180.0', '1110.0', '190.0', '1080.0', '1090.0', '1220.0', '1170.0', '1500.0', '160.0', '1140.0', '170.0', '490.0', '1180.0', '1150.0', '210.0', '1160.0', '130.0', '1280.0', '1320.0', '90.0', '1260.0', '1380.0', '1240.0', '1330.0', '80.0', '1360.0', '1340.0', '1290.0', '1420.0', '1390.0', '1600.0', '1350.0', '1460.0', '1310.0', '1590.0', '1430.0', '1580.0', '1440.0', '1510.0', '1540.0', '230.0', '60.0', '1480.0', '1490.0', '1650.0', '1780.0', '1690.0', '1760.0', '1570.0', '1720.0', '1520.0', '1620.0', '1870.0', '1530.0', '1790.0', '1680.0', '70.0', '1850.0', '1940.0', '1550.0', '1470.0', '1710.0', '2020.0', '1640.0', '1830.0', '1900.0', '1630.0', '1950.0', '40.0', '1610.0', '1860.0', '2160.0', '1750.0', '2170.0', '2070.0', '2150.0', '265.0', '414.0', '1810.0', '2330.0', '1840.0', '2000.0', '2010.0', '2040.0', '515.0', '2100.0', '2030.0', '2080.0', '1820.0', '2550.0', '435.0', '1890.0', '235.0', '2090.0', '2110.0', '2190.0', '2610.0', '1008.0', '946.0', '666.0', '1245.0', '1525.0', '1880.0', '862.0', '2300.0', '768.0', '274.0', '20.0', '2810.0', '508.0', '143.0', '417.0', '556.0', '915.0', '207.0', '295.0', '2120.0', '2310.0', '266.0', '1275.0', '225.0', '176.0', '516.0', '602.0', '1248.0', '276.0', '2180.0', '1990.0', '1548.0', '2506.0', '588.0', '2850.0', '1284.0', '875.0', '2570.0', '2500.0', '3000.0', '2490.0', '4130.0', '1481.0', '11852.0', '2360.0', '2600.0', '243.0', '704.0', '784.0', '2390.0', '374.0', '518.0', '935.0', '792.0', '475.0', '1930.0', '2196.0', '652.0', '415.0', '3260.0', '1913.0', '4820.0', '2050.0', '1960.0', '1920.0', '3480.0']
0.0      12826
?         454
600.0     217
500.0     209
700.0     208
...
1920.0      1
3480.0      1
2730.0      1
2720.0      1
248.0       1
Name: sqft_basement, Length: 304, dtype: int64
```

From the output above, we can see that this is numeric data that has been converted to a string because of the presence of the '?' character. In order to make this data usable, we shall convert it to a float. Furthermore, we will need to deal with the missing values ('?') in this column. The missing values represent 2.1% of the total records in the dataset. As this is a small percentage of the total records, we shall be dropping the records missing values.

#### 2.1.2.10 Yr Built

The `yr_built` column identifies the year the house was built.

```
In [295]:
```


```
# Describe the 'yr_built' column
```

```
describe_data(data, 'yr_built')
```

```
# Visualise the data distribution
```

```
plot_distribution(data, 'yr_built', 'Year Built Column Data Distribution', 115)
```

```
count    21597.000000
mean      1970.999676
std        29.375234
min       1900.000000
25%       1951.000000
50%       1975.000000
75%       1997.000000
max       2015.000000
Name: yr_built, dtype: float64
```



From the distributions above we can see that the data is slightly skewed to the left. This is because the mean is less than the median. The oldest house in the dataset was built in 1900, and the newest house in the dataset was built in 2015. The mean year the houses in the dataset were built is 1971, and the median year the houses were built is 1975. The standard deviation of the yr built column is 29.

### 2.1.2.11 Yr Renovated

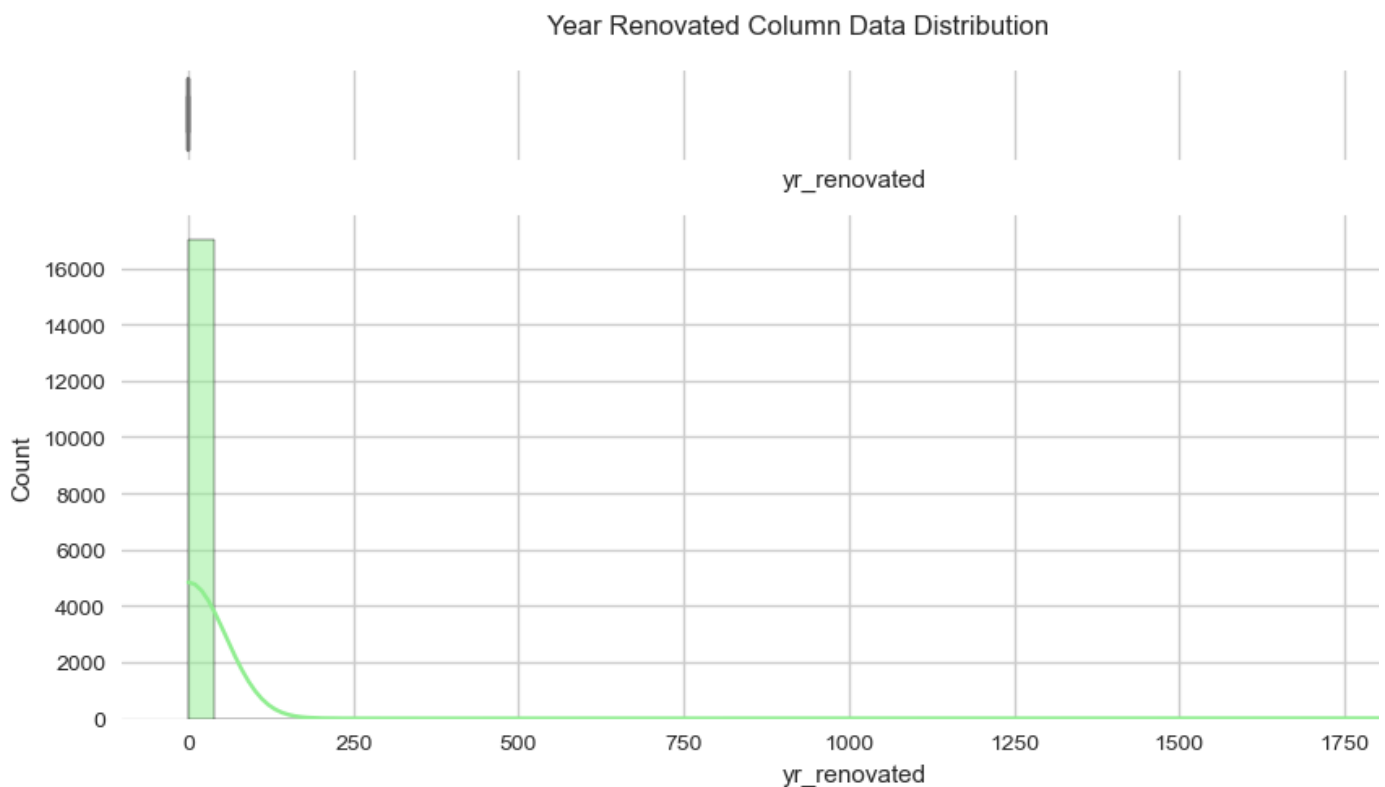
The `yr_renovated` column identifies the year the house was renovated.

In [296]:

```
# Describe the 'yr_renovated' column
describe_data(data, 'yr_renovated')

# Visualise the data distribution
plot_distribution(data, 'yr_renovated', 'Year Renovated Column Data Distribution', 50)
```

```
count    17755.000000
mean       83.636778
std       399.946414
min         0.000000
25%         0.000000
50%         0.000000
75%         0.000000
max      2015.000000
Name: yr_renovated, dtype: float64
```



From the distribution and value counts above, we can see that the data has a number of zero values, suggesting that the house has not been renovated, or that the data is missing. Furthermore, there are missing values in this column. We shall be analysing the data more in depth in the next phase to see how to handle missing values in the column.

### 2.1.2.12 Lat & Long

The `lat` column identifies the latitude of the house. The `long` column identifies the longitude of the house.

In [297]:

```
latlon = list(zip(data.lat, data.long))

base_map = folium.Map([data.lat.mean(), data.long.mean()], zoom_start=13)
base_map

for coord in latlon:
    folium.Marker(location=[coord[0], coord[1]], fill_color='#43d9de', radius=8).add_to(base_map)

# export the map as HTML file
base_map.save('../images/map.html')
```

From the exported map above, we can see that the houses in the dataset are located in the same area. Importantly, we see that the houses are roughly within the same area therefore we do not need to provide additional data in the `lat` and `long` columns.

### 2.1.2.13 Sqft Living15

The `sqft_living15` square footage of interior housing living space for the nearest 15 neighbors.

```
In [298]:
```

```
# Describe the 'sqft_living15' column
```


```
describe_data(data, 'sqft_living15')
```

```
# Visualise the data distribution
```

```
plot_distribution(data, 'sqft_living15', 'Sqft Living15 Column Data Distribution', 100)
```

```
count    21597.000000
mean      1986.620318
std       685.230472
min       399.000000
25%      1490.000000
50%      1840.000000
75%      2360.000000
max       6210.000000
```

```
Name: sqft_living15, dtype: float64
```



From the distributions above, we can see that the data is skewed to the right. This is as a result of the mean being greater than the median. The minimum square footage of the nearest 15 neighbors is 399, and the maximum square footage of the nearest 15 neighbors is 6,210. The mean square footage of the nearest 15 neighbors is 1987. The standard deviation of the nearest 15 neighbors is 1840. The standard deviation of the sqft living15 column is 1840.

#### 2.1.2.14 Sqft Lot15

---

The `sqft_lot15` column represents the square footage of the land lots for the nearest 15 neighbors.

---

```
In [299]:
```


```
# Describe the 'sqft_lot15' column
```

```
describe_data(data, 'sqft_lot15')
```

```
# Visualise the data distribution
```

```
plot_distribution(data, 'sqft_lot15', 'Sqft Lot15 Column Data Distribution', 100)
```

```
count      21597.000000
mean       12758.283512
std        27274.441950
min         651.000000
25%        5100.000000
50%        7620.000000
75%       10083.000000
max       871200.000000
Name: sqft_lot15, dtype: float64
```



In the distributions above we see a much more skewed to the right column. The minimum sq neighbors is 651, and the maximum square footage of the nearest 15 neighbors is 871,200. nearest 15 neighbors is 12758, and the median square footage of the nearest 15 neighbors i of the sqft lot15 column is 27274.

## Summary Of Numerical Columns

- The data in the numerical columns is also of fairly decent quality. Other than a few missi column and datatype corrections that need to be made to the date and sqft basement c



There are quite a number of outliers in the data, however, I do not think that will affect the

## 3. Data Processing

This phase, which is often referred to as “data munging”, prepares the final data set(s) and involves the following tasks:

- Select Data
- Clean Data
- Construct Data
- Integrate Data
- Format Data

### 3.1 Clean Data

In this section we will be looking at the missing values in the dataset as well as the duplicates in the dataset.

The columns that were identified to be having missing data and duplicates were:

- id
- waterfront
- yr renovated
- view

#### 3.1.1 Duplicate Records

The `id` column was identified to have duplicate records. However, we did not know if the duplicates were the same records or if they were different records with the same id. In order to find out, we shall be looking at the records in the `id` column.

```
In [300]:
```

```
# Create a new dataframe that contains the ids that have been duplicated in the dataset
duplicates = data[data.duplicated(['id'], keep=False)]

# Preview the duplicates dataframe
duplicates
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
93	6021501535	7/25/2014	430000.0	3	1.50	1580	5000	1.0	NO
94	6021501535	12/23/2014	700000.0	3	1.50	1580	5000	1.0	NO
313	4139480200	6/18/2014	1380000.0	4	3.25	4290	12103	1.0	NO
314	4139480200	12/9/2014	1400000.0	4	3.25	4290	12103	1.0	NO
324	7520000520	9/5/2014	232000.0	2	1.00	1240	12092	1.0	NaN
...	...	...	...	...	...	...	...	...	...
20654	8564860270	3/30/2015	502000.0	4	2.50	2680	5539	2.0	NaN
20763	6300000226	6/26/2014	240000.0	4	1.00	1200	2171	1.5	NO
20764	6300000226	5/4/2015	380000.0	4	1.00	1200	2171	1.5	NO
21564	7853420110	10/3/2014	594866.0	3	3.00	2780	6000	2.0	NO
21565	7853420110	5/4/2015	625000.0	3	3.00	2780	6000	2.0	NO

353 rows × 10 columns

Looking at the duplicated id records, we can see that the records are not erroneous. The same house was sold multiple times. Therefore, we shall be keeping the records. In order to check for duplicate records, we shall be checking the date column along with the id to see if there are any duplicate records on the same day. That would be an erroneous record.

```
In [301]:
```

```
# Check for duplicate records that have both the same id and date
duplicates[duplicates.duplicated(['id', 'date'], keep=False)]
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_basement
--	----	------	-------	----------	-----------	-------------	----------	--------	------------	------	-----	-------	---------------

0 rows × 14 columns

We see that we have no records indicating that the same house has been sold multiple times; we shall be keeping the records with duplicated ids.

### 3.1.2 Missing Values

The columns that were identified to be having missing data were `waterfront`, `yr_renovated` with the missing values in these columns. Furthermore, using the insights that were identified, we shall be using the type of data along with the data distribution to determine the best way to deal with the missing values.

#### 3.1.2.1 Waterfront

The `waterfront` column is a categorical column. The column has 2 unique values, 'YES' and 'NO'. 'YES' accounted for 11% of the total records in the dataset. As this is a fairly large percentage, replacing the missing values with the mode of the column. The mode of the column is 'NO'. We will replace the missing values with 'NO'.

```
In [302]:
```

```
# Fill the missing values with the mode of the column  
data['waterfront'] = data['waterfront'].fillna(data['waterfront'].mode()[0])
```

#### 3.1.2.2 Year Renovated

The `yr_renovated` column is a numerical column. With 3842 missing values, this accounted for 11% of the total records in the dataset. Furthermore, majority of the data in the records were zero. This could either be because the house has never been renovated or that the data is erroneous. As there is no ideal way of dealing with this, we will drop the entire column.

```
In [303]:
```

```
# Drop the 'yr_renovated' column
data.drop('yr_renovated', axis=1, inplace=True)

# Preview the first five rows of the dataframe
data.head()
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	NO	NC
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	NO	NC
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	NO	NC
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	NO	NC
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	NO	NC

### 3.1.2.3 View

The `view` column is a categorical column. With 63 missing values, this accounted for 0.3% of the total records. As this is a small percentage of the total records, we shall be dropping the records with missing values.

```
In [304]:
```

```
# Drop the missing records in the 'view' column
data = data[data.view.notnull()]

get_value_counts(data, 'view')
```

```
NONE      19422
AVERAGE    957
GOOD       508
FAIR       330
EXCELLENT  317
Name: view, dtype: int64
```

## 3.2 Construct Data

In this section, we shall be deriving new attributes that will be helpful in our analysis.

We shall be creating new columns that will be useful in the analysis. The columns that we sh and price columns which represents the date that the homes were sold and their price resp from the date column and creating a new column called `yr_sold` . The `yr_sold` column will b to adjust the price of the homes for inflation, if so, the column will once again be used to calc

### 3.2.1 Year Sold

---

The `yr_sold` column represents the year that the homes were sold.

---


```
In [305]:
```

```
# Create a new column, 'yr_sold', from the 'date' column
data['yr_sold'] = data['date'].apply(lambda x: int(x.split('/')[1]))

# View the values (and counts) in the 'yr_sold' column
print(get_value_counts(data, 'yr_sold'))

# Visualise the data distribution
plot_data(data, 'yr_sold', 'Year Sold Column Data Distribution')
```

```
2014    14588
2015     6946
Name: yr_sold, dtype: int64
```



We see that the `yr_sold` column has been created and populated with the year that the homes in the dataset were sold in 2014. Ultimately, the data is fairly clean and good to w

### 3.2.2 Current Price

The `current_price` column represents the price of the homes adjusted for inflation.


The current price column could be created since there are different years involved in the sales. The prices of the homes may be different due to the different market conditions. Therefore, by creating the current price of the homes, we can fairly compare the prices of the homes in different years. To create this column, we first have to establish that the prices of the homes are indeed different due to the different market conditions. We shall be doing this by looking at the distributions of the prices of the homes sold in the different years.

In [306]:

```
# Create different dataframes for each year (2014 and 2015)
df_2014 = data[data['yr_sold'] == 2014]
df_2015 = data[data['yr_sold'] == 2015]

# Plot the distribution of the 'price' column for each year
fig, ax = plt.subplots(figsize=(10, 5))

plt.hist(df_2014['price'], bins=100, color='orange', alpha=0.7, label='House Prices in 2014')
plt.hist(df_2015['price'], bins=100, color='lightblue', alpha=0.7, label='House Prices in 2015')
plt.title('House Prices Over The Years')
plt.legend();
```



From the distribution above, we can see that there is not much difference in the prices of the

years. Therefore, there is no need to create a new column with the current price of the home sold column as it is no longer needed.

```
In [307]:  
  
# Drop the 'yr_sold' column  
data.drop('yr_sold', axis=1, inplace=True)
```

### 3.3 Format Data

In this section, we shall be re-formatting data as necessary.

The specific columns that we shall be looking at in this section are:

- date
- sqft\_basement

#### 3.3.1 Date

We shall be converting the date column to a datetime object.

```
In [308]:  
  
# Convert the 'date' column to datetime format  
data['date'] = pd.to_datetime(data['date'], format='%m/%d/%Y')  
  
# Preview the first five rows of the dataframe  
data.head()
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
0	7129300520	2014-10-13	221900.0	3	1.00	1180	5650	1.0	NO	NONE
1	64141400192	2014-12-09	538000.0	3	2.25	2570	7242	2.0	NO	NONE
2	5631500400	2015-02-25	180000.0	2	1.00	770	10000	1.0	NO	NONE
3	2487200875	2014-12-09	604000.0	4	3.00	1960	5000	1.0	NO	NONE
4	1954400510	2015-02-18	510000.0	3	2.00	1680	8080	1.0	NO	NONE

#### 3.3.2 Basement Square Footage



We shall be converting the `sqft_basement` column to a numerical column. However, we first r  
'?' value that we identified in the previous phase. In the previous phase we saw that it accou  
in the dataset. As this is a fairly small percentage of the total records, we shall be dropping th  
value. This will ensure that we are not introducing any bias into the dataset. Once we have d  
erroneous '?' value, we shall be converting the remaining `sqft_basement` column values to a

```
In [309]:  
  
# Drop the records with a '?' in the 'sqft_basement' column  
data = data[data['sqft_basement'] != '?']  
  
# Convert the 'sqft_basement' column to float  
data['sqft_basement'] = data['sqft_basement'].astype(float)  
  
# Preview the first five rows of the dataframe  
data.head()
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
0	7129300520	2014-10-13	221900.0	3	1.00	1180	5650	1.0	NO	NONE
1	6414100192	2014-12-09	538000.0	3	2.25	2570	7242	2.0	NO	NONE
2	5631500400	2015-02-25	180000.0	2	1.00	770	10000	1.0	NO	NONE
3	2487200875	2014-12-09	604000.0	4	3.00	1960	5000	1.0	NO	NONE
4	1954400510	2015-02-18	510000.0	3	2.00	1680	8080	1.0	NO	NONE

```
In [310]:
```

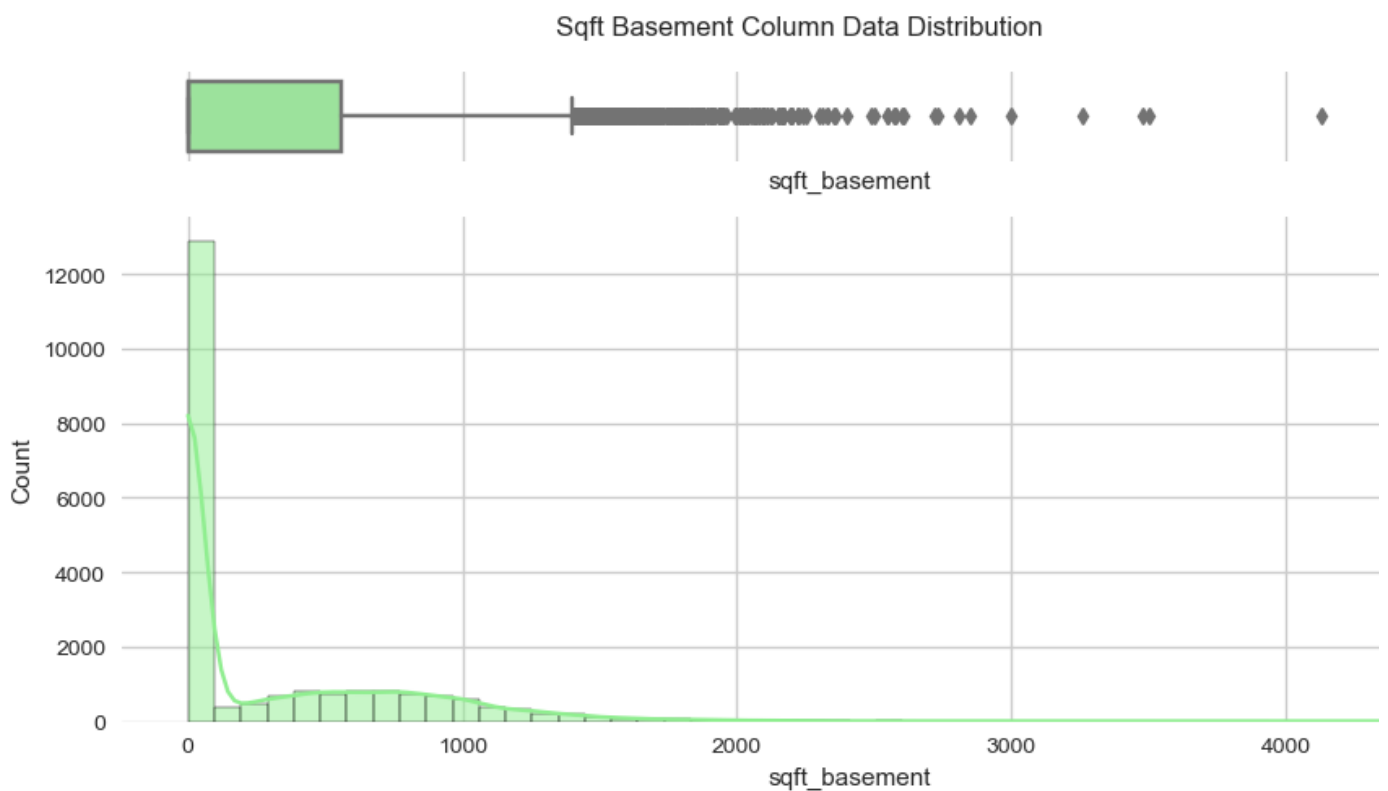
```
# We can also visualize the data distribution of the 'sqft_basement' column
```

```
describe_data(data, 'sqft_basement')
```

```
# Plot the visualisation
```

```
plot_distribution(data, 'sqft_basement', 'Sqft Basement Column Data Distribution', 50)
```

```
count    21082.000000
mean      291.359975
std       442.007858
min        0.000000
25%        0.000000
50%        0.000000
75%       560.000000
max      4820.000000
Name: sqft_basement, dtype: float64
```



From our distributions above, we can see that the 'sqft\_basement' column is highly positively skewed, the mean is higher than the median. This is as a result of the outliers in the data. Furthermore, the maximum basement size is 4820 square feet. This is quite a large basement size. However, as it is, it is not erroneous.

Now that we have completely cleaned our data, we can export the cleaned data to a csv file.

```
In [311]:
```

```
# Export the dataframe to a csv file
data.to_csv('../data/processed/cleaned_kc_house_data.csv', index=False)
```

## 4. Modeling

In this phase, we'll likely build and assess various models based on several different features. This phase has four tasks:

- Select Modeling Techniques
- Generate Test Design
- Build Models
- Assess Models

### 4.1 Select Modeling Techniques

In this section, we shall be determining which algorithms to try

I believe that the best algorithm to try for this experiment is regression. Regression is a supervised learning algorithm used to predict the value of a dependent variable based on the value of the independent variables. By using regression to estimate the effect that the different features of the homes have on our dependent variable, the price, the result will be able to provide our stakeholder with a model that will be able to predict the price of the homes that will have the most impact on the price of the homes.

Furthermore, as we are working with multiple features, we will be using multiple linear regression. Multiple linear regression is a regression algorithm that is used to predict the value of a dependent variable based on multiple independent variables (unlike linear regression which only uses one independent variable).

### 4.2 Build Models

---

In this section, we shall be building the models.

---

We will first start by building a baseline model. The baseline model will be used to compare the models that we will be building. After that, we will build our multiple linear regression model.

### 4.2.1 Build Baseline Simple Linear Regression Model

---

A baseline model is essentially a simple model that acts as a reference in a machine function is to contextualize the results of trained models.

---

The target variable is price. Therefore, we look at the correlation coefficients for all of the features with the highest correlation with price.

In [312]:

```
# Create a correlation matrix for the dataset
corr = data.corr()['price'].sort_values(ascending=False)
corr
```

```
price          1.000000
sqft_living    0.702004
sqft_above     0.605481
sqft_living15  0.586495
bathrooms      0.525029
sqft_basement  0.323018
bedrooms       0.308454
lat            0.307667
floors         0.256603
sqft_lot       0.088400
sqft_lot15     0.083530
yr_built       0.054849
long           0.022512
id            -0.016413
zipcode        -0.053429
Name: price, dtype: float64
```

We see that the `sqft_living` column has the highest correlation with the `price` column. This indicates that the size of the house is a major factor in determining the price of the house. We shall also create a scatter plot between the `sqft_living` and `price`.

```
In [313]:
```

```
# Plot a scatter plot of the 'price' column against the 'sqft_living' column
```

```
plt.figure(figsize=(10, 5))
```

```
plt.scatter(data['sqft_living'], data['price'], color='lightgreen', alpha=0.7, s=10, ec
```

```
plt.title('Price vs Living Space')
```

```
plt.xlabel('Living Space (sqft)')
```

```
plt.ylabel('Price');
```



We can now declare `y` and `X_baseline` variables, where `y` is a Series containing price data and `X_baseline` is a DataFrame containing the column with the strongest correlation (`sqft_living`).

```
In [314]:
```

```
# Declare y and X_baseline variables
y = data['price']
X_baseline = data[['sqft_living']]
```

Next, we'll use our variables to build and fit a simple linear regression model

```
In [315]:
```

```
# Create a baseline model
baseline_model = sm.OLS(y, sm.add_constant(X_baseline))
baseline_results = baseline_model.fit()

# Print the summary results of the baseline model
print(baseline_results.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.493
Model:                  OLS       Adj. R-squared:            0.493
Method:                 Least Squares   F-statistic:          2.048e+04
Date:                  Sat, 01 Oct 2022   Prob (F-statistic):    0.00
Time:                  11:55:31    Log-Likelihood:       -2.9287e+05
No. Observations:      21082      AIC:                  5.857e+05
Df Residuals:          21080      BIC:                  5.858e+05
Df Model:               1
Covariance Type:       nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const          -4.327e+04    4456.393     -9.709     0.000    -5.2e+04    -3.45e+04
sqft_living     280.4877       1.960    143.116     0.000     276.646     284.329
=====
Omnibus:            14303.984    Durbin-Watson:           1.986
Prob(Omnibus):      0.000      Jarque-Bera (JB):       509767.330
Skew:               2.786      Prob(JB):               0.00
Kurtosis:           26.437      Cond. No.               5.63e+03
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.63e+03. This might indicate that there are strong multicollinearity or other numerical problems.

We can plot the regression line on top of the scatter plot earlier to see how well the model fits

In [316]:

```
# Plot a scatter plot of the 'price' column against the 'sqft_living' column
```

```
plt.figure(figsize=(10, 5))
```

```
# Plot the regression line of the baseline model
```

```
x = np.linspace(data.sqft_living.min(), data.sqft_living.max(), 100)
```

```
Y = baseline_results.params[0] + baseline_results.params[1] * x
```

```
plt.plot(x, Y, color='black', label='Regression Line')
```

```
plt.scatter(data['sqft_living'], data['price'], color='lightgreen', alpha=0.7, s=10, ec
```

```
plt.title('Price vs Living Space (Baseline Model)')
```

```
plt.xlabel('Living Space (sqft)')
```

```
plt.ylabel('Price (\$)')
```

```
plt.legend();
```



```
In [317]:
```

```
# Calculate the mean absolute error of the baseline model
baseline_mae = mean_absolute_error(y, baseline_results.predict(sm.add_constant(X_base):
baseline_mae
```

```
173713.2378046139
```

Our most strongly correlated variable with price is sqft\_living

The model is statistically significant as it explains only 50% of the variance in the data. However, our analysis. In a typical prediction, the model is off by about \$173992.

- The intercept is about -45130. *This means that if a zero square foot house would be built, the predicted price would be about -\$45,130.*
- The coefficient of sqft\_living is about 281. *This means that for every square foot increase in the house, the price of the house increases by about \$281.*

#### 4.2.2 Build Iterated Multiple Linear Regression Model

We will now iterate the baseline model by building a multiple linear regression model with more than one independent variable.

We will start by creating a new dataframe that will contain all of the features that we want to use in our model. We will also encode the categorical columns. In order to know which variables to keep in our model, we will create a correlation matrix. This is done in order to reduce multicollinearity. Multicollinearity is a situation in which two or more variables are highly correlated. This can cause problems in the model as it can lead to unstable coefficient estimates. Therefore, we will be removing the variables that are highly correlated with each other.



```
In [318]:
```

```
# Declare X_iter variables
```

```
X_iter = data[['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront'
```

```
# Preview the X_iter dataframe
```

```
X_iter
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft
0	3	1.00	1180	5650	1.0	NO	NONE	Average	7 Average	1180
1	3	2.25	2570	7242	2.0	NO	NONE	Average	7 Average	2170
2	2	1.00	770	10000	1.0	NO	NONE	Average	6 Low Average	770
3	4	3.00	1960	5000	1.0	NO	NONE	Very Good	7 Average	1050
4	3	2.00	1680	8080	1.0	NO	NONE	Average	8 Good	1680
...	...	...	...	...	...	...	...	...	...	...
21592	3	2.50	1530	1131	3.0	NO	NONE	Average	8 Good	1530
21593	4	2.50	2310	5813	2.0	NO	NONE	Average	8 Good	2310
21594	2	0.75	1020	1350	2.0	NO	NONE	Average	7 Average	1020
21595	3	2.50	1600	2388	2.0	NO	NONE	Average	8 Good	1600
21596	2	0.75	1020	1076	2.0	NO	NONE	Average	7 Average	1020

21082 rows x 14 columns

We have 4 categorical columns in our dataset. As a result, we will need to encode them in our model. We will be ordinal encoding the `condition` and `grade` columns and one-hot encode the `view` and `waterfront` columns.

#### 4.2.2.1 Encode Categorical Columns

We will now encode the categorical columns in the dataset.

##### 4.2.2.1.1 Ordinal Encoding

Ordinal encoding converts each label into integer values and the encoded data represents the labels.

Using the official [King County Assessor Website \(https://info.kingcounty.gov/assessor/esales\)](https://info.kingcounty.gov/assessor/esales) we were able to understand that the values in the `condition` and `grade` columns are ordinal, an based on the quality of the feature. Therefore, we will be ordinal encoding these columns.

```
In [319]:  
  
# Create dictionaries for mapping the ordinal numerical value  
condition_dict = {'Poor': 1, 'Fair': 2, 'Average': 3, 'Good': 4, 'Very Good': 5}  
grade_dict = {'3 Poor': 3, '4 Low': 4, '5 Fair': 5, '6 Low Average': 6, '7 Average': 7,  
  
# Map the ordinal numerical values to the 'condition' and 'grade' columns  
X_iter['condition'] = X_iter['condition'].map(condition_dict)  
X_iter['grade'] = X_iter['grade'].map(grade_dict)  
  
# Preview the dataframe  
X_iter
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_
0	3	1.00	1180	5650	1.0	NO	NONE	3	7	1180
1	3	2.25	2570	7242	2.0	NO	NONE	3	7	2170
2	2	1.00	770	10000	1.0	NO	NONE	3	6	770
3	4	3.00	1960	5000	1.0	NO	NONE	5	7	1050
4	3	2.00	1680	8080	1.0	NO	NONE	3	8	1680
...	...	...	...	...	...	...	...	...	...	...
21592	3	2.50	1530	1131	3.0	NO	NONE	3	8	1530
21593	4	2.50	2310	5813	2.0	NO	NONE	3	8	2310
21594	2	0.75	1020	1350	2.0	NO	NONE	3	7	1020
21595	3	2.50	1600	2388	2.0	NO	NONE	3	8	1600
21596	2	0.75	1020	1076	2.0	NO	NONE	3	7	1020

21082 rows x 14 columns

### 4.2.2.1.2 One Hot Encoding

One hot encoding is a process of converting categorical data variables so they can be used by machine learning algorithms to improve predictions.

We shall be encoding the remaining categorical columns (waterfront and view ) using one hot encoding to avoid the "Dummy Variable Trap" (perfect multicollinearity between the independent

In [320]:

```
# Encode the categorical variables
X_iter = pd.get_dummies(X_iter, columns=['waterfront', 'view'], drop_first=False)

# Preview the dataframe
X_iter
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	grade	sqft_above	sqft_basem
0	3	1.00	1180	5650	1.0	3	7	1180	0.0
1	3	2.25	2570	7242	2.0	3	7	2170	400.0
2	2	1.00	770	10000	1.0	3	6	770	0.0
3	4	3.00	1960	5000	1.0	5	7	1050	910.0
4	3	2.00	1680	8080	1.0	3	8	1680	0.0
...	...	...	...	...	...	...	...	...	...
21592	3	2.50	1530	1131	3.0	3	8	1530	0.0
21593	4	2.50	2310	5813	2.0	3	8	2310	0.0
21594	2	0.75	1020	1350	2.0	3	7	1020	0.0
21595	3	2.50	1600	2388	2.0	3	8	1600	0.0
21596	2	0.75	1020	1076	2.0	3	7	1020	0.0

21082 rows x 19 columns

In the waterfront column, we shall be dropping the waterfront\_NO column as the reference study the effect of having a house on a waterfront. In the view column, we shall be dropping reference column. This will allow us to study the effect of having a house with a view. In addition, we shall be dropping the view\_NO column as the reference value in the column.

```
In [321]:
```

```
# Drop the 'waterfront_NO' and 'view_NONE' columns
X_iter.drop(['waterfront_NO', 'view_NONE'], axis=1, inplace=True)

# Preview the dataframe
X_iter
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	grade	sqft_above	sqft_basem
0	3	1.00	1180	5650	1.0	3	7	1180	0.0
1	3	2.25	2570	7242	2.0	3	7	2170	400.0
2	2	1.00	770	10000	1.0	3	6	770	0.0
3	4	3.00	1960	5000	1.0	5	7	1050	910.0
4	3	2.00	1680	8080	1.0	3	8	1680	0.0
...	...	...	...	...	...	...	...	...	...
21592	3	2.50	1530	1131	3.0	3	8	1530	0.0
21593	4	2.50	2310	5813	2.0	3	8	2310	0.0
21594	2	0.75	1020	1350	2.0	3	7	1020	0.0
21595	3	2.50	1600	2388	2.0	3	8	1600	0.0
21596	2	0.75	1020	1076	2.0	3	7	1020	0.0

21082 rows x 17 columns

#### 4.2.2.2 Correlation Matrix

A correlation matrix is a table showing correlation coefficients between variables

We will now analyse the correlation matrix to determine which variables to keep in our model. We will also be looking at the VIF (Variance Inflation Factor) of each variable. The VIF variance of an estimated regression coefficient increases if the independent variables are correlated.

We are aiming to ensure that a correlation coefficient is less than 0.6 and a VIF is less than 5. A correlation coefficient of 0.6 or higher indicates that the variables are highly correlated. A VIF of 5 or higher indicates that the variables are highly correlated.

In [322]:

*# Define function to plot the correlation matrix*

```
def corrmatrix(df):  
    ''' This function plots a correlation matrix for a given dataframe '''  
    plt.figure(figsize=(10, 5))  
  
    corr = df.corr()  
  
    # Generate a mask to only show the bottom triangle  
    mask = np.triu(np.ones_like(corr, dtype=bool))  
  
    # generate heatmap  
    sns.heatmap(round(corr,2), annot=True, mask=mask, vmin=-1, vmax=1, cmap='Greens')  
    plt.title('Correlation Coefficient Of Predictors')  
    plt.show()  
  
# Define function to print the VIF values of the predictors  
def vif_df(df):  
    ''' This function prints the VIF values of the predictors in a given dataframe '''  
    vif_data = pd.DataFrame()  
  
    # Add a constant to the dataframe  
    X = df.assign(const=1)  
  
    vif_data["feature"] = X.columns  
  
    # calculating VIF for each feature  
    vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]  
  
    print(vif_data.sort_values(by='VIF', ascending=False))  
  
# Plot the correlation matrix  
corrmatrix(X_iter)  
  
# Print the VIF values of the predictors  
vif_df(X_iter)
```

Correlation Coefficient Of Predictors

bedrooms																		
bathrooms	0.51																	
sqft_living	0.58	0.75																
sqft_lot	0.03	0.09	0.17															
floors	0.18	0.5	0.35	-0.01														
condition	0.03	-0.13	-0.06	-0.01	-0.26													
grade	0.36	0.67	0.76	0.11	0.46	-0.15												
sqft_above	0.48	0.69	0.88	0.18	0.52	-0.16	0.76											
sqft_basement	0.3	0.28	0.43	0.02	-0.25	0.17	0.17	-0.05										
yr_built	0.16	0.51	0.32	0.05	0.49	-0.36	0.45	0.43	-0.13									
sqft_living15	0.39	0.57	0.76	0.14	0.28	-0.09	0.71	0.73	0.2	0.33								
sqft_lot15	0.03	0.09	0.18	0.72	-0.01	-0	0.12	0.2	0.02	0.07	0.18							
waterfront_YES	-0	0.06	0.1	0.02	0.02	0.02	0.08	0.07	0.08	-0.02	0.08	0.03						
view_AVERAGE	0.05	0.09	0.13	0.04	0.01	0.03	0.12	0.08	0.13	-0.05	0.14	0.04	0					
view_EXCELLENT	0.04	0.11	0.17	0.02	0.03	0.03	0.15	0.11	0.15	-0.02	0.15	0.03	0.56	-0.03				
view_FAIR	0.02	0.04	0.07	-0.01	-0.02	0.02	0.05	0.02	0.1	-0.03	0.08	-0.01	-0.01	-0.03	-0.02			
view_GOOD	0.05	0.11	0.16	0.07	0.02	0.02	0.14	0.09	0.16	-0.02	0.16	0.06	0.04	-0.03	-0.02	-0.02		
	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	grade	sqft_above	sqft_basement	yr_built	sqft_living15	sqft_lot15	waterfront_YES	view_AVERAGE	view_EXCELLENT	view_FAIR	view_GOOD	

	feature	VIF
2	sqft_living	inf
7	sqft_above	inf
8	sqft_basement	inf
17	const	7991.177201
1	bathrooms	3.302484
6	grade	3.240348
10	sqft_living15	2.812048
11	sqft_lot15	2.123483
3	sqft_lot	2.094685
4	floors	1.934228
9	yr_built	1.816720
0	bedrooms	1.641691
14	view_EXCELLENT	1.546534
12	waterfront_YES	1.478656
5	condition	1.187949
16	view_GOOD	1.079328
13	view_AVERAGE	1.063058
15	view_FAIR	1.025156

We can see a very high VIF value for the sqft\_living , sqft\_above , and sqft\_basement colour

dropping `sqft_living` from our model.

```
In [323]:
```

```
# Drop the 'sqft_living' column
```


```
X_iter.drop(['sqft_living'], axis=1, inplace=True)
```

```
# Visualize the correlation matrix and the VIF dataframe
```

```
corrmatrix(X_iter)
```

```
vif_df(X_iter)
```

Correlation Coefficient Of Predictors



	feature	VIF
16	const	7991.177201
6	sqft_above	4.834813
1	bathrooms	3.302484
5	grade	3.240348
9	sqft_living15	2.812048
10	sqft_lot15	2.123483
2	sqft_lot	2.094685
7	sqft_basement	1.981861
3	floors	1.934228
8	yr_built	1.816720
0	bedrooms	1.641691
13	view_EXCELLENT	1.546534
11	waterfront_YES	1.478656



4	condition	1.187949
15	view_GOOD	1.079328
12	view_AVERAGE	1.063058
14	view_FAIR	1.025156

The correlation matrix shows that the `sqft_above` column still has a high correlation. Therefore, we should remove `sqft_above` from our model.

```
In [324]:
```

```
# Drop the 'sqft_above' column
```

```
X_iter.drop(['sqft_above'], axis=1, inplace=True)
```

```
# Visualize the correlation matrix and the VIF dataframe
```

```
corrmatrix(X_iter)
```

```
vif_df(X_iter)
```

Correlation Coefficient Of Predictors

bedrooms															
bathrooms	0.51														
sqft_lot	0.03	0.09													
floors	0.18	0.5	-0.01												
condition	0.03	-0.13	-0.01	-0.26											
grade	0.36	0.67	0.11	0.46	-0.15										
sqft_basement	0.3	0.28	0.02	-0.25	0.17	0.17									
yr_built	0.16	0.51	0.05	0.49	-0.36	0.45	-0.13								
sqft_living15	0.39	0.57	0.14	0.28	-0.09	0.71	0.2	0.33							
sqft_lot15	0.03	0.09	0.72	-0.01	-0	0.12	0.02	0.07	0.18						
waterfront_YES	-0	0.06	0.02	0.02	0.02	0.08	0.08	-0.02	0.08	0.03					
view_AVERAGE	0.05	0.09	0.04	0.01	0.03	0.12	0.13	-0.05	0.14	0.04	0				
view_EXCELLENT	0.04	0.11	0.02	0.03	0.03	0.15	0.15	-0.02	0.15	0.03	0.56	-0.03			
view_FAIR	0.02	0.04	-0.01	-0.02	0.02	0.05	0.1	-0.03	0.08	-0.01	-0.01	-0.03	-0.02		
view_GOOD	0.05	0.11	0.07	0.02	0.02	0.14	0.16	-0.02	0.16	0.06	0.04	-0.03	-0.02	-0.02	
	bedrooms	bathrooms	sqft_lot	floors	condition	grade	sqft_basement	yr_built	sqft_living15	sqft_lot15	waterfront_YES	view_AVERAGE	view_EXCELLENT	view_FAIR	view_GOOD

	feature	VIF
15	const	7956.618704
1	bathrooms	2.969652
5	grade	2.780924
8	sqft_living15	2.282297
9	sqft_lot15	2.114176
2	sqft_lot	2.082816
3	floors	1.896715
7	yr_built	1.790623
6	sqft_basement	1.578408
12	view_EXCELLENT	1.546500
10	waterfront_YES	1.477030
0	bedrooms	1.475156
4	condition	1.187719

14	view_GOOD	1.078949
11	view_AVERAGE	1.062833
13	view_FAIR	1.025087

The correlation coefficients for the `bathrooms` column is still higher than our threshold. There `bathrooms` column from our model.

```
In [325]:
```

```
# Drop the 'bathrooms' column
```


```
X_iter.drop(['bathrooms'], axis=1, inplace=True)
```

```
# Visualize the correlation matrix and the VIF dataframe
```

```
corrmatrix(X_iter)
```

```
vif_df(X_iter)
```

Correlation Coefficient Of Predictors



	feature	VIF
14	const	7066.323721
4	grade	2.621639
7	sqft_living15	2.249374
8	sqft_lot15	2.114127
1	sqft_lot	2.080635
2	floors	1.670595
6	yr_built	1.615203
11	view_EXCELLENT	1.546476
9	waterfront_YES	1.476606
5	sqft_basement	1.401570
0	bedrooms	1.323164
3	condition	1.185202
13	view_GOOD	1.078677

10	view_AVERAGE	1.062599
12	view_FAIR	1.025008

Dropping the `bathrooms` column has further reduced the overall correlation in the dataset. However, `sqft_lot15` columns still have a high correlation. Therefore, we will be dropping them both from the dataset.

```
In [326]:
```

```
# Drop the 'sqft_lot15' and 'sqft_living15' column
```


```
X_iter.drop(['sqft_lot15', 'sqft_living15'], axis=1, inplace=True)
```

```
# Visualize the correlation matrix and the VIF dataframe
```

```
corrmatrix(X_iter)
```

```
vif_df(X_iter)
```

Correlation Coefficient Of Predictors



	feature	VIF
12	const	6995.603649
4	grade	1.753385
2	floors	1.658031
6	yr_built	1.606184
9	view_EXCELLENT	1.537416
7	waterfront_YES	1.476263
5	sqft_basement	1.400994
0	bedrooms	1.266129
3	condition	1.184633
11	view_GOOD	1.068646
8	view_AVERAGE	1.054079
1	sqft_lot	1.024057
10	view_FAIR	1.019902

Now that we have our VIF and correlation matrix, below the threshold, we can now build our model.

#### 4.2.2.3 Build Model

---

We will now build our multiple linear regression model.

---

In [327]:

```
# Create a model
iterated_model = sm.OLS(y, sm.add_constant(X_iter))
iterated_results = iterated_model.fit()

# Print the summary results of the baseline model
print(iterated_results.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.601
Model:                  OLS       Adj. R-squared:            0.601
Method:                 Least Squares   F-statistic:            2647.
Date:                   Sat, 01 Oct 2022   Prob (F-statistic):      0.00
Time:                   11:55:37   Log-Likelihood:         -2.9033e+05
No. Observations:       21082   AIC:                    5.807e+05
Df Residuals:           21069   BIC:                    5.808e+05
Df Model:                12
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	5.137e+06	1.33e+05	38.501	0.000	4.88e+06	5.4e+06
bedrooms	1.624e+04	1940.787	8.368	0.000	1.24e+04	2e+04
sqft_lot	0.1365	0.039	3.482	0.000	0.060	0.213
floors	7.826e+04	3808.798	20.548	0.000	7.08e+04	8.57e+04
condition	1.944e+04	2669.062	7.284	0.000	1.42e+04	2.47e+04
grade	2.051e+05	1799.964	113.942	0.000	2.02e+05	2.09e+05
sqft_basement	118.8650	4.272	27.822	0.000	110.491	127.239
yr_built	-3277.1706	68.955	-47.526	0.000	-3412.328	-3142.013
waterfront_YES	5.409e+05	2.38e+04	22.745	0.000	4.94e+05	5.88e+05
view_AVERAGE	6.994e+04	7976.568	8.768	0.000	5.43e+04	8.56e+04
view_EXCELLENT	3.409e+05	1.64e+04	20.776	0.000	3.09e+05	3.73e+05
view_FAIR	1.311e+05	1.3e+04	10.052	0.000	1.06e+05	1.57e+05
view_GOOD	1.371e+05	1.09e+04	12.595	0.000	1.16e+05	1.58e+05

```

=====
Omnibus:                18594.679   Durbin-Watson:            1.974
Prob(Omnibus):           0.000   Jarque-Bera (JB):         2185504.878
Skew:                    3.741   Prob(JB):                  0.00
Kurtosis:                52.315   Cond. No.                  3.67e+06
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.67e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Now we can compare the baseline and iterated model statistics.



In [328]:

```
# Compare the baseline and iterated model mean absolute errors
iterated_mae = mean_absolute_error(y, iterated_results.predict(sm.add_constant(X_iter)))
print("Baseline Model Mean Absolute Error: ", baseline_mae)
print("Iterated Model Mean Absolute Error: ", iterated_mae)

# Compare the adjusted R-squared values of the baseline and model
print("Baseline Model Adjusted R-squared: ", baseline_results.rsquared_adj)
print("Iterated Model Adjusted R-squared: ", iterated_results.rsquared_adj)
```

```
Baseline Model Mean Absolute Error: 173713.2378046139
Iterated Model Mean Absolute Error: 147296.625677394
Baseline Model Adjusted R-squared: 0.49278505895823355
Iterated Model Adjusted R-squared: 0.6010210227347927
```

From the model results, we can see that the model is statistically significant as it explains 60 compared to the 49% in the baseline model. Furthermore, the model is off by about 147, 29' baseline model. This is a significant improvement.

We will now do an analysis of the coefficients of the model.

```
In [329]:
```

```
# Create a dataframe of the coefficients of the iterated model along with their p-values
results_df = pd.concat([round(iterated_results.params,3), round(iterated_results.pvalues,3)], axis=1)
results_df.columns = ["coefficient", "p-value"]
results_df
```

	coefficient	p-value
const	5137454.742	0.000000
bedrooms	16239.726	0.000000
ssqft_lot	0.137	0.000498
floors	78262.473	0.000000
condition	19442.583	0.000000
grade	205091.669	0.000000
ssqft_basement	118.865	0.000000
yr_built	-3277.171	0.000000
waterfront_YES	540924.431	0.000000
view_AVERAGE	69940.897	0.000000
view_EXCELLENT	340899.246	0.000000
view_FAIR	131058.754	0.000000
view_GOOD	137050.033	0.000000

#### 4.2.2.4 Model Results Analysis

We can see that all of the variables in the model are statistically significant.

- We can see that constant value is about 5,137,455. This means that a house with no features would be worth about 5,137,455.
- The coefficient of bedrooms is 16,240 which means that for every bedroom increase in the house, the price of the house increases by about 16,240.
- The coefficient of ssqft\_lot is 0.14 which means that for every square foot increase in the lot, the price of the house increases by about 0.14.
- The coefficient of floors is 78,262 which means that for every floor increase in the house, the price of the house increases by about 78,262.
- The coefficient of condition is 19,443 which means that for every condition rating increase in the house, the price of the house increases by about 19,443.

- The coefficient of `grade` is  
*205,092 which means that for every grade rating increase in the house, the price of the house increases by about 205,092.*
- The coefficient of `sqft_basement` is  
*119 which means that for every square foot increase in the basement, the price of the house increases by about 119.*
- The coefficient of `yr_built` is -  
*3,277 which means that for every year increase in the year the house was built, the price of the house decreases by about 3,277.*
- The coefficient of `waterfront_YES` is  
*540,924 this means that if a house is on a water front, the price of the house increases 540,924.*
- The coefficients for `view` range from 69,941 to 340,899
  - `view_AVERAGE` is  
*69,941 which means that for an average view compared to no view, the price of the house increases by about 69,941.*
  - `view_FAIR` is  
*131,058 which means that for a fair view compared to no view, the price of the house increases by about 131,058.*
  - `view_GOOD` is  
*137,050 which means that for a good view compared to no view, the price of the house increases by about 137,050.*
  - `view_EXCELLENT` is  
*340,899 which means that for an excellent view compared to no view, the price of the house increases by about 340,899.*
- This `view` outcome is surprising since we would expect that the effect of having an average view is better than having a fair view. However, the model shows that the effect of having a fair view is better than having an average view. This could perhaps suggest that the homes in the dataset with an average view are not as well valued as those with a fair view, or that those homes with an average value have been undervalued.

## 5. Conclusion

In this phase we will be interpreting the model results and limitations in the context of and giving recommendations to the stakeholder based on our modeling results.

## 5.1 Recommendations

Taking this analysis back to the original business problem, the aim was to help a real estate agent find the best possible potential renovations to make to increase the value. After modelling the renovations are as follows:

- Moving the house closer to the water. This will increase the value of the house by about 540,924. As a result this most likely means that it makes the view excellent as the two features. In turn, by making the view excellent, the value of the house will increase by about 340,899. However, this renovation can only be made if land is close water.
- The second best renovation to make is to improve the grade of the house. This will increase about \$205,092 for every grade.
- The third best renovation to make is to increase the number of floors in the house. This will increase the value of the house by about \$78,262 for every floor. However, it is worth mentioning that our data only had 1 floor as the maximum. Therefore, it is unlikely that this statistic would apply to a house with more than 1 floor.
- Increasing the number of bedrooms in the house will increase the value of the house by about \$15,000 per bedroom. However, it is worth mentioning that our data only had 10 bedrooms as the maximum. Therefore, it is unlikely that this statistic would apply to a house with more than 10 bedrooms.
- Lastly, increasing the size of the basement will increase the value of the house by about \$10,000 per square foot. However, it is worth mentioning that our data only had 4,820 square feet as the maximum. Therefore, it is unlikely that this statistic would apply to a house with more than 4,820 square feet.

## 5.2 Limitations

Though our model did show a significant increase in the accuracy of the model, there are still limitations. These limitations are as follows:

- The data in the dataset is from 2014 and 2015. Therefore, it may not be able to account for changes in the real estate market since then. As a result the model may not be able to predict the value of a house accurately.
- In order to improve the value of a house, we would need to understand the market (i.e. what buyers are looking for). Therefore, by not having this information, we are unable to advise our clients on the best possible renovation to build the most expensive house in the world, but if it is not what buyers are looking for, there is no value in that.
- By using a correlation threshold of 0.6, we may have ignored dropping some features which may have led to multicollinearity in the model. As a result, the model may not be able to predict the value of a house accurately.